

Your Title Here

Courtney Guo

Massachusetts Institute of
Technology
ckguo@mit.edu

Danielle Pace

Massachusetts Institute of
Technology
dfpace@csail.mit.edu

Polina Golland

Massachusetts Institute of
Technology
polina@csail.mit.edu

ABSTRACT

Congenital heart disease (CHD) affects many children around the world, which sometimes require extensive surgery planning. For doctors to be able to plan surgeries, scans of the heart must be accurately labelled with the different substructures of the heart. This project aims to develop methods to localize internal structures in the heart in MRI scans of patients with congenital heart disease. Localization will be done by applying a regression forest to the MRI scans of the patient, to find the bounding boxes of the structures within the heart. A regression forest is a supervised machine learning model that is a collection of regression trees, which is described later in the technical approach.

1 INTRODUCTION

Millions of children around the world are born with congenital heart defects, which may require surgery to be treated. For doctors to effectively plan surgery, it is useful if a model of the heart can be 3D-printed, which requires the scans of the patient to be segmented. Traditionally, this means that an expert must manually label each voxel in the image with whether or not it belongs in the bloodpool, myocardium, or outside the heart [correct me if I'm wrong]. Currently, patient-specific 3D heart models are underused because it takes around 4-8 hours to manually segment cardiac MRI images, since each contains approximately 1503 voxels. There have been algorithms developed to segment the heart in normal adult patients, one of the most popular being atlas segmentation, which uses a fully segmented heart as a reference and tries to adapt a patient's heart to the given reference. However, atlas-based methods do not work well on children with congenital heart disease (CHD) due to the irregular location or shapes of the organs.

This paper proposes an alternative method to aid the segmentation of hearts with CHD. Segmenting images of hearts with CHD as opposed to healthy hearts is an additional challenge, because hearts with CHD may have incomplete or missing structures, or structures located in different areas compared to healthy hearts. Therefore, many traditional methods for segmenting these images do not work well. The method proposed in this paper attempts to address these challenges. A first step to segmentation is locating bounding boxes around the regions of interest, as there exists algorithms that will produce accurate segmentation given a bounding box of the region to be labeled.

The method proposed in this paper will localize the heart as well as structures within the heart in patients with congenital heart disease, using random forests.

2 RELATED WORK

One method for segmentation of the heart in children with CHD is an interactive algorithm proposed by Danielle Pace. At each step of the algorithm, the user is directed to manually label one slice of the heart that will give the most information [1]. The algorithm then segments each target slice according to its closest reference slices. To segment a patch, the algorithm finds the k most similar patches in the set of relevant reference regions, and “similarity” depends on patch intensities, gradients, and positions, and each pixel is labeled according to a majority vote. This algorithm greatly reduces the amount of time needed to segment a heart, but still requires user interaction.

Regression forests have also been shown to do well in image segmentation, specifically the detection and localization of organs. A. Criminisi has applied regression forests to learn the non-linear mapping from voxels directly to organ position and size [2]. The regression forest groups voxels with similar features or similar field of views together, and learns an estimate of the bounding boxes of each organ using the training data that reaches that node. Intuitively, each voxel contributes varying degrees of confidence to the estimates of the location and size of every organ. When applied on real datasets, the forest learns to recognize key indicators (such tips of the ribs or vertebrae) and those pixels provide high confidence estimates of where certain organs are located. Criminisi then compared these results against other methods, including Elastix

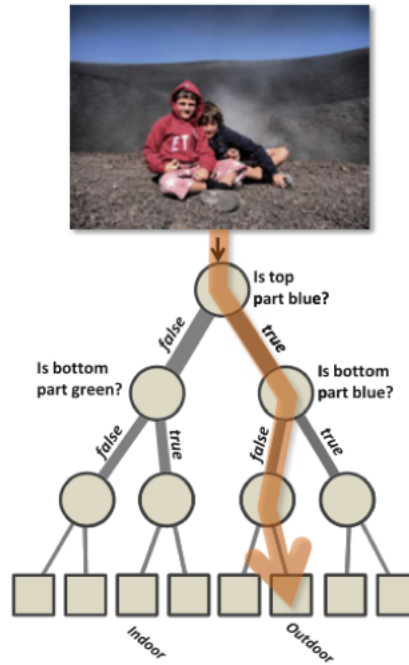


Figure 1: A decision tree

and Simplex methods, as well as atlas methods, and showed that the regression forest method was superior in accuracy. This paper extends the regression forest methods used by Criminisi to children with CHD.

3 METHOD

3.1 Regression Random Forests

This paper develops a method to localize the heart in 3D MRI scans of patients with CHD, based on regression random forests. A random forest is a supervised machine learning model that consists of many decision trees, in which the training phase of each tree contains some amount of randomness. A decision tree is a flow-chart like structure in which each node is a Boolean function on the data's features. For each data point that passes through the decision tree, as it arrives at each node, the node makes a decision for whether the data point goes to its left child or right child, based on the data point's features. Figure 1 depicts a decision tree that takes in an image as input and predicts whether it was taken indoors or outdoors.

This paper deals with regression forests, which contain regression trees instead of decision trees. Regression trees are the same as decision trees except each leaf node contains a regression model instead of a




Figure 2: A regression forest




Figure 3: Training a regression forest

classification model. Therefore each leaf node would predict a regression value instead of a classification label.

A regression forest is produced by aggregating many regression trees. Figure 2 shows a regression forest consisting of only 3 trees, and for each tree a data point can be seen passing through the tree's branches and down to the leaf node, where it is put into a Gaussian distribution.

A regression forest can be trained by individually training its regression trees. Regression trees are trained by minimizing some loss function at each node. This paper uses information gain, given below:

$$I(S, \theta) = H(S) - \frac{|S^L|}{|S|}H(S^L) - \frac{|S^R|}{|S|}H(S^R) \quad (1)$$

where S is the dataset that reaches that node, θ denotes the parameters of the function being considered at the split node, S^L is the subset that will go to the left node, and S^R is the subset that will go to the right node, and $H(S)$ is the entropy of a dataset S , given below:

$$H(S) = \frac{1}{2}n \log(2\pi e|\Sigma|) \quad (2)$$

where Σ is the covariance matrix of the Gaussian model that is fitted to the dataset at the node.

To train a regression tree, we greedily choose the split functions at each node. At each node, K features are randomly selected, and the feature that maximizes the information gain given above is chosen as the split function for that node. Figure 3 depicts how a selected feature divides the training set, and then Gaussian distributions are fitted to the resulting split nodes. After each regression tree is trained, the set of these trees becomes the regression forest. The introduction of randomness in feature selection is meant to prevent overfitting: each tree is a weak learner, but the aggregation of many randomly trained weak learners should be a strong learner. Once a regression forest is trained, it can make predictions on new data points by running the data point through each of the regression trees until it reaches a leaf node, and combining the outputs of each regression tree.

There are three hyper-parameters that determine the structure of the regression forest: T is the number of trees, D is the maximum tree depth, and K is the number of features being considered at each node.

3.2 Regression Forests for Localization of the Heart

We trained a regression forest to predict the location of the bounding box of the heart given MRI scans of patients with congenital heart disease. A bounding box is given by a 6-dimensional vector: $(x_1, x_2, y_1, y_2, z_1, z_2)$. x_1 and x_2 denote the x-values of the two faces of the bounding box that are perpendicular to the x-axis, where x_1 is the smaller of the two values, and y_1, y_2, z_1, z_2 are similarly defined for the other axes. Then, for

each voxel that we run through the regression forest, we calculate its offset from the bounding box:

$$d(\mathbf{p}, \mathbf{b}) = (x_1, x_2, y_1, y_2, z_1, z_2) - (p_x, p_x, p_y, p_y, p_z) \quad (3)$$

The Gaussian model fitted at each node is then a 6D Gaussian, representing the voxel’s predicted offset from the bounding box.

To produce predictions for the bounding box of a heart given a 3D image, we take each voxel in the image and run it through the regression forest. For each voxel, we can combine the Gaussian distributions resulting from each tree in the forest by sampling the distributions. Then, we add the voxel’s location to get the distribution of the predicted absolute location of the bounding box. Then, we combine the distributions from each voxel to form the predicted distribution of the absolute bounding box of the heart.

$$\hat{\mathbf{b}} = \operatorname{argmax}_b \left(\sum_{i=0}^{|S|} \sum_{t=0}^T p(\text{something??}) \right) \quad (4)$$

3.3 Feature Selection

As mentioned previously, at each split node, K features are chosen randomly for consideration as the split feature. Each feature is calculated as the mean intensity of voxels in a rectangular box, that is at some offset to the voxel. Therefore, each feature is governed by the following parameters: θ is the offset to the center of the box, which is a 3D vector, and ϕ is the size of the box, which is also a 3D vector. Each dimension of the box must be odd, to ensure that the center of the box is on a lattice point.

$$F_{\theta, \phi} = \frac{1}{|B|} \sum_{v \in B} J(v) \quad (5)$$

where B is the set of voxels that lie in the box that is described by an offset of θ and a box size of ϕ , and $J(v)$ describes the intensity at the voxel v . These types of features were also used by Criminisi et al.

To generate a random feature, we sample θ and ϕ from uniform distributions. Each element of θ is sampled from a uniform distribution from 0 to a fixed fraction of the size of the image in that dimension. Each element of ϕ is sampled from a uniform distribution of 1 to $2k + 1$ for some k . We have not been able to tune parameters much, and currently the parameters are set at $\frac{1}{6}$ of the image size for θ , and $k = 5$ for ϕ .

4 THE BODY OF THE PAPER

4.1 Paper Structure

We recommend using the following section titles to structure your paper. Abstract, Introduction, Related Works, Details on Technical Work, Implementation, Evaluation, Results, and Conclusion. Keep in mind that this may change for different communities and types of works. Based on your review of other publications in your area, edit this document accordingly.

4.2 Format

Typically, the body of a paper is organized into a hierarchical structure, with numbered or unnumbered headings for sections, subsections, sub-subsections, and even smaller sections. The command `\section` that precedes this paragraph is part of such a hierarchy.¹ L^AT_EX handles the numbering and placement of these headings for you, when you use the appropriate heading commands around the titles of the headings. If you want a sub-subsection or smaller part to be unnumbered in your output, simply append an asterisk to the command name. Examples of both numbered and unnumbered headings will appear throughout the balance of this sample document.

Because the entire article is contained in the **document** environment, you can indicate the start of a new paragraph with a blank line in your input file; that is why this sentence forms a separate paragraph.

4.3 Type Changes and *Special* Characters

We have already seen several typeface changes in this sample. You can indicate italicized words or phrases in your text with the command `\textit`; emboldening with the command `\textbf` and typewriter-style (for instance, for computer code) with `\texttt`. But remember, you do not have to indicate typestyle changes when such changes are part of the *structural* elements of your article; for instance, the heading of this subsection will be in a sans serif² typeface, but that is handled by the document class file. Take care

¹This is a footnote.

²Another footnote here. Let's make this a rather long one to see how it looks.

with the use of³ the curly braces in typeface changes; they mark the beginning and end of the text that is to be in the different typeface.

You can use whatever symbols, accented characters, or non-English characters you need anywhere in your document; you can find a complete list of what is available in the *L^AT_EX User's Guide* [?].

4.4 Math Equations

You may want to display math equations in three distinct styles: inline, numbered or non-numbered display. Each of the three are discussed in the next sections.

4.4.1 Inline (In-text) Equations. A formula that appears in the running text is called an inline or in-text formula. It is produced by the **math** environment, which can be invoked with the usual `\begin . . . \end` construction or with the short form `$. . . $`. You can use any of the symbols and structures, from α to ω , available in L^AT_EX [?]; this section will simply show a few examples of in-text equations in context. Notice how this equation: $\lim_{n \rightarrow \infty} x = 0$, set here in in-line math style, looks slightly different when set in display style. (See next section).

4.4.2 Display Equations. A numbered display equation — one set off by vertical space from the text and centered horizontally — is produced by the **equation** environment. An unnumbered display equation is produced by the **displaymath** environment.

Again, in either environment, you can use any of the symbols and structures available in L^AT_EX; this section will just give a couple of examples of display equations in context. First, consider the equation, shown as an inline equation above:

$$\lim_{n \rightarrow \infty} x = 0 \tag{6}$$

Notice how it is formatted somewhat differently in the **displaymath** environment. Now, we'll enter an unnumbered equation:

$$\sum_{i=0}^{\infty} x + 1$$

³Another footnote.

and follow it with another numbered equation:

$$\sum_{i=0}^{\infty} x_i = \int_0^{\pi+2} f \tag{7}$$

just to demonstrate L^AT_EX’s able handling of numbering.

4.5 Citations

Citations to articles [? ? ? ?], conference proceedings [?] or maybe books [? ?] listed in the Bibliography section of your article will occur throughout the text of your article. You should use BibTeX to automatically produce this bibliography; you simply need to insert one of several citation commands with a key of the item cited in the proper location in the .tex file [?]. The key is a short reference you invent to uniquely identify each work; in this sample document, the key is the first author’s surname and a word from the title. This identifying key is included with each item in the .bib file for your article.

The details of the construction of the .bib file are beyond the scope of this sample document, but more information can be found in the *Author’s Guide*, and exhaustive details in the *L^AT_EX User’s Guide* by Lamport [?].

This article shows only the plainest form of the citation command, using `\cite`.

4.6 Tables

Because tables cannot be split across pages, the best placement for them is typically the top of the page nearest their initial cite. To ensure this proper “floating” placement of tables, use the environment **table** to enclose the table’s contents and the table caption. The contents of the table itself must go in the **tabular** environment, to be aligned properly in rows and columns, with the desired horizontal and vertical rules. Again, detailed instructions on **tabular** material are found in the *L^AT_EX User’s Guide*.

Immediately following this sentence is the point at which Table 1 is included in the input file; compare the placement of the table here with the table in the printed output of this document.

To set a wider table, which takes up the whole width of the page’s live area, use the environment **table*** to enclose the table’s contents and the table caption. As with a single-column table, this wide table will “float”

Table 1: Frequency of Special Characters

Non-English or Math	Frequency	Comments
Ø	1 in 1,000	For Swedish names
π	1 in 5	Common in math
\$	4 in 5	Used in business
Ψ^2_1	1 in 40,000	Unexplained usage

Table 2: Some Typical Commands

Command	A Number	Comments
<code>\author</code>	100	Author
<code>\table</code>	300	For tables
<code>\table*</code>	400	For wider tables

to a location deemed more desirable. Immediately following this sentence is the point at which Table 2 is included in the input file; again, it is instructive to compare the placement of the table here with the table in the printed output of this document.

It is strongly recommended to use the package `booktabs` [?] and follow its main principles of typography with respect to tables:

- (1) Never, ever use vertical rules.
- (2) Never use double rules.

It is also a good idea not to overuse horizontal rules.

4.7 Figures

Like tables, figures cannot be split across pages; the best placement for them is typically the top or the bottom of the page nearest their initial cite. To ensure this proper “floating” placement of figures, use the environment **figure** to enclose the figure and its caption.

This sample document contains examples of `.eps` files to be displayable with \LaTeX . If you work with pdf\LaTeX , use files in the `.pdf` format. Note that most modern \TeX systems will convert `.eps` to `.pdf` for you on the fly. More details on each of these are found in the *Author’s Guide*.

Figure 4: A sample black and white graphic that needs to span two columns of text.

Figure 5: A sample black and white graphic that has been resized with the `includegraphics` command.

As was the case with tables, you may want a figure that spans two columns. To do this, and still to ensure proper “floating” placement of tables, use the environment **figure*** to enclose the figure and its caption. And don’t forget to end the environment with **figure***, not **figure**!

4.8 Theorem-like Constructs

Other common constructs that may occur in your article are the forms for logical constructs like theorems, axioms, corollaries and proofs. ACM uses two types of these constructs: theorem-like and definition-like.

Here is a theorem:

THEOREM 4.1. Let f be continuous on $[a, b]$. If G is an antiderivative for f on $[a, b]$, then

$$\int_a^b f(t) dt = G(b) - G(a).$$

Here is a definition:

Definition 4.2. If z is irrational, then by e^z we mean the unique number that has logarithm z :

$$\log e^z = z.$$

The pre-defined theorem-like constructs are **theorem**, **conjecture**, **proposition**, **lemma** and **corollary**. The pre-defined definition-like constructs are **example** and **definition**. You can add your own constructs using the *amsthm* interface [?]. The styles used in the `\theoremstyle` command are **acmplain** and **acmdefinition**.

Another construct is **proof**, for example,

PROOF. Suppose on the contrary there exists a real number L such that

$$\lim_{x \rightarrow \infty} \frac{f(x)}{g(x)} = L.$$

Then

$$l = \lim_{x \rightarrow c} f(x) = \lim_{x \rightarrow c} \left[g(x) \cdot \frac{f(x)}{g(x)} \right] = \lim_{x \rightarrow c} g(x) \cdot \lim_{x \rightarrow c} \frac{f(x)}{g(x)} = 0 \cdot L = 0,$$

which contradicts our assumption that $l \neq 0$. □

5 CONCLUSIONS

This paragraph will end the body of this sample document. Remember that you might still have Acknowledgments or Appendices; brief samples of these follow. There is still the Bibliography to deal with; and we will make a disclaimer about that here: with the exception of the reference to the L^AT_EX book, the citations in this paper are to articles which have nothing to do with the present subject and are used as examples only.

A HEADINGS IN APPENDICES

The rules about hierarchical headings discussed above for the body of the article are different in the appendices. In the **appendix** environment, the command **section** is used to indicate the start of each Appendix, with alphabetic order designation (i.e., the first is A, the second B, etc.) and a title (if you include one). So, if you need hierarchical structure *within* an Appendix, start with **subsection** as the highest level. Here is an outline of the body of this document in Appendix-appropriate form:

A.1 Introduction

A.2 The Body of the Paper

A.2.1 Type Changes and Special Characters.

A.2.2 Math Equations.

Inline (In-text) Equations.

Display Equations.

A.2.3 Citations.

A.2.4 Tables.

A.2.5 Figures.

A.2.6 Theorem-like Constructs.

A Caveat for the T_EX Expert.

A.3 Conclusions

A.4 References

Generated by bibtex from your `.bib` file. Run latex, then bibtex, then latex twice (to resolve references) to create the `.bbl` file. Insert that `.bbl` file into the `.tex` source file and comment out the command `\thebibliography`.

B MORE HELP FOR THE HARDY

Of course, reading the source code is always useful. The file `acmart.pdf` contains both the user guide and the commented code.