

Localization of the Heart in MRI Scans

Courtney Guo

Massachusetts Institute of
Technology
ckguo@mit.edu

Danielle Pace

Massachusetts Institute of
Technology
dfpace@csail.mit.edu

Polina Golland

Massachusetts Institute of
Technology
polina@csail.mit.edu

ABSTRACT

Congenital heart disease (CHD) affects many children around the world, which sometimes require extensive surgery planning. For doctors to be able to plan surgeries, it is helpful if scans of the heart are accurately labelled with the different substructures of the heart. This project aims to develop methods to localize internal structures in the heart as well as the heart itself in MRI scans of patients with congenital heart disease. Localization will be done by applying a regression forest to the MRI scans of the patient, to find the bounding boxes of the structures within the heart. A regression forest is a supervised machine learning model that is a collection of regression trees, which is described later in the technical approach.

1 INTRODUCTION

Millions of children around the world are born with congenital heart defects, which may require surgery to be treated. For doctors to effectively plan surgery, it is useful if a model of the heart can be 3D-printed, which requires the scans of the patient to be segmented. Traditionally, this means that an expert must manually label each voxel in the image with whether or not it belongs in the bloodpool, myocardium, or outside the heart. Currently, patient-specific 3D heart models are underused because it takes around 4-8 hours to manually segment cardiac MRI images, since each contains 100-150 slices. There have been algorithms developed to segment the heart in normal adult patients, one of the most popular being atlas segmentation, which uses a fully segmented heart as a reference and tries to align a patient's heart to the given reference. However, atlas-based methods do not work well on children with congenital heart disease (CHD) due to the irregular location or shapes of the organs.

This paper proposes an alternative method to aid the segmentation of hearts with CHD. Segmenting images of hearts with CHD as opposed to healthy hearts is an additional challenge, because hearts with CHD may have incomplete or missing structures, or structures located in different areas compared to healthy hearts. Therefore, many traditional methods for segmenting these images do not work well. The method proposed in this paper attempts to address these challenges. A first step to segmentation is locating bounding boxes around the regions of interest, which will assist downstream processing.

The method proposed in this paper will localize the heart in patients with congenital heart disease, using random forests.

2 RELATED WORK

One method for segmentation of the heart in children with CHD is an interactive algorithm proposed by Danielle Pace. The user is directed to manually label 10-15 slices uniformly distributed throughout the 3D volume [1]. The algorithm then segments each remaining slice according to its closest reference slices. To segment a patch, the algorithm finds the k most similar patches in the set of relevant reference regions, and "similarity" depends on patch intensities, gradients, and positions, and each pixel is labeled according to a majority vote. This algorithm greatly reduces the amount of time needed to segment a heart, but still requires user interaction.

Regression forests have also been shown to do well in image segmentation, specifically the detection and localization of organs. A. Criminisi has applied regression forests to learn the non-linear mapping from voxels directly to organ position and size [2]. The regression forest groups voxels with similar features or similar field of views together, and learns an estimate of the bounding boxes of each organ using the training data that reaches that node. Intuitively, each voxel contributes varying degrees of confidence to the estimates of the location and size of every organ. When applied on real datasets, the forest learns to recognize key indicators (such tips of the ribs or vertebrae) and those pixels provide high confidence estimates of where certain organs are located. Criminisi then compared these results against other methods, including Elastix and Simplex methods, as well as atlas methods, and showed that the regression forest method was superior in accuracy. This paper applies the regression forest methods used by Criminisi to children with CHD.

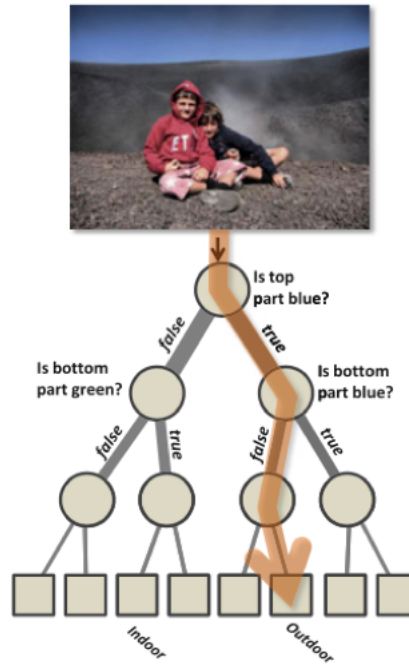


Figure 1: A decision tree

3 METHOD

3.1 Regression Random Forests

This paper develops a method to localize the heart in 3D MRI scans of patients with CHD, based on regression random forests. A random forest is a supervised machine learning model that consists of many decision trees, in which the training phase of each tree contains some amount of randomness. A decision tree is a flow-chart like structure in which each node is a Boolean function on the data's features. For each data point that passes through the decision tree, as it arrives at each node, the node makes a decision for whether the data point goes to its left child or right child, based on the data point's features. At each leaf node, the data point is classified with a label, which can be binary or multi-class classification. Figure 1 depicts a decision tree that takes in an image as input and predicts whether it was taken indoors or outdoors.

This paper deals with regression forests, which contain regression trees instead of decision trees. Regression trees are the same as decision trees except each leaf node contains a regression model instead of a

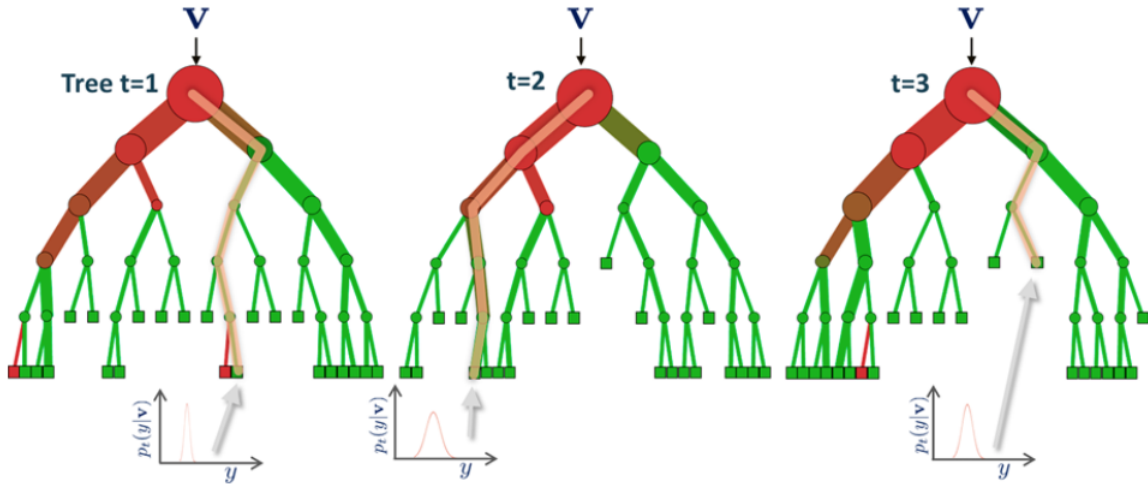


Figure 2: A regression forest

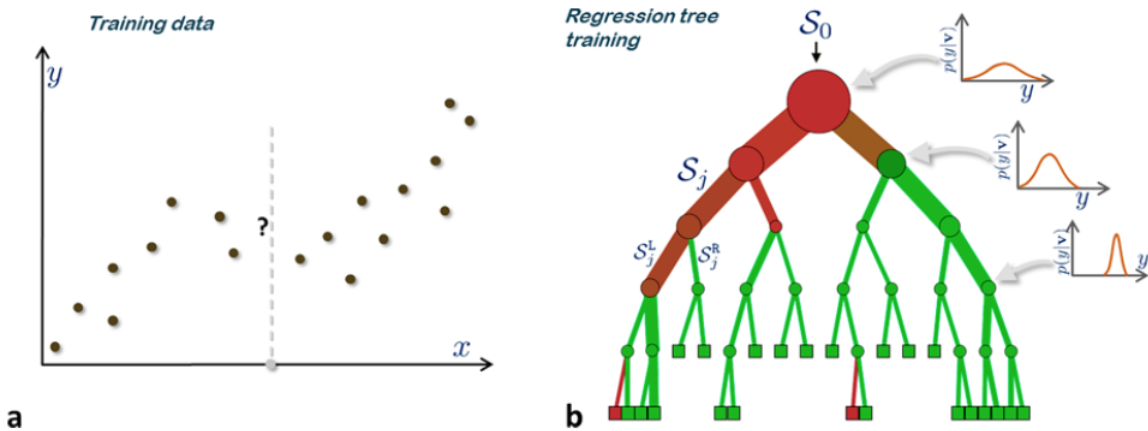


Figure 3: Training a regression forest

classification model. Therefore each leaf node would predict a regression value instead of a classification label.

A regression forest is produced by aggregating many regression trees. Figure 2 shows a regression forest consisting of only 3 trees, and for each tree a data point can be seen passing through the tree's branches and down to the leaf node, where it is put into a Gaussian distribution.

A regression forest can be trained by individually training its regression trees. Regression trees are trained by minimizing some loss function at each node. This paper uses information gain, given below:

$$I(S, \theta) = H(S) - \frac{|S^L|}{|S|}H(S^L) - \frac{|S^R|}{|S|}H(S^R) \quad (1)$$

where S is the dataset that reaches that node, θ denotes the parameters of the function being considered at the split node, S^L is the subset that will go to the left node, and S^R is the subset that will go to the right node, and $H(S)$ is the entropy of a dataset S , given below:

$$H(S) = \frac{1}{2}n \log(2\pi e|\Sigma|) \quad (2)$$

where Σ is the covariance matrix of the Gaussian model that is fitted to the dataset at the node. We maximize information gain as the method for choosing a split function, because it will choose the function that will separate the dataset the most. Intuitively, this results in greedily training the tree such that each time the tree branches, we get the most information out of the dataset that ends up at a certain branch.

To train a regression tree, we greedily choose the split functions at each node. At each node, K features are randomly selected, and the feature that maximizes the information gain given above is chosen as the split function for that node. Figure 3 depicts how a selected feature divides the training set, and then Gaussian distributions are fitted to the resulting split nodes. After each regression tree is trained, the set of these trees becomes the regression forest. The introduction of randomness in feature selection is meant to prevent overfitting: each tree is a weak learner, but the aggregation of many randomly trained weak learners should be a strong learner. Training many weak learners in parallel is also computationally much faster than training one strong learner such as a deep regression tree. Once a regression forest is trained, it can make predictions on new data points by running the data point through each of the regression trees until it reaches a leaf node, and combining the outputs of each regression tree. Training a tree terminates either when the maximum tree depth has been reached, or when the information gain is less than some threshold value.

There are three hyper-parameters that determine the structure of the regression forest: T is the number of trees, D is the maximum tree depth, and K is the number of features being considered at each node.

3.2 Regression Forests for Localization of the Heart

We trained a regression forest to predict the location of the bounding box of the heart given MRI scans of patients with congenital heart disease. A bounding box is given by a 6-dimensional vector: $(x_1, x_2, y_1, y_2, z_1, z_2)$. x_1 and x_2 denote the x-values of the two faces of the bounding box that are perpendicular to the x-axis, where x_1 is the smaller of the two values, and y_1, y_2, z_1, z_2 are similarly defined for the other axes. Instead of running each entire image through the tree, we run features for each voxel through the tree, and each of those will define a distribution of where the bounding box is. We can aggregate the distributions from each voxel to get a final distribution for where the bounding box should be.

For each voxel $\mathbf{p} = (p_x, p_y, p_z)$ that we run through the regression forest, we calculate its offset from the bounding box \mathbf{b} :

$$d(\mathbf{p}, \mathbf{b}) = (x_1, x_2, y_1, y_2, z_1, z_2) - (p_x, p_x, p_y, p_y, p_z, p_z) \quad (3)$$

The Gaussian model fitted at each node is then a 6D Gaussian, representing the voxel's predicted offset from the bounding box. Each of the leaf nodes has a distribution of the voxel's offset from the bounding box, not the bounding box locations themselves.

To produce predictions for the bounding box of a heart given a 3D image, we take each voxel in the image and run it through the regression forest. For each voxel, we can combine the Gaussian distributions resulting from each tree in the forest by sampling the distributions. Then, we add the voxel's location to get the distribution of the predicted absolute location of the bounding box. Then, we combine the distributions from each voxel to form the predicted distribution of the absolute bounding box of the heart.

3.3 Feature Selection

As mentioned previously, at each split node, K features are chosen randomly for consideration as the split feature. Each feature is calculated as the mean intensity of voxels in a rectangular box, that is at some offset to the voxel. Therefore, each feature is governed by the following parameters: θ is the offset to the center of the box, which is a 3D vector, and ϕ is the size of the box, which is also a 3D vector. Each dimension of the box must be odd, to ensure that the center of the box is on a lattice point.

$$F_{\theta,\phi} = \frac{1}{|B|} \sum_{v \in B} J(v) \quad (4)$$

where B is the set of voxels that lie in the box that is described by an offset of θ and a box size of ϕ , and $J(v)$ describes the intensity at the voxel v . These types of features were also used by Criminisi et al. Presence of certain features in the image, such as bright or dark spots, can hint at where the heart is. Therefore, having an especially high or low mean intensity of voxels at a specific offset can contribute a good prediction as to where the heart may be.

To generate a random feature, we sample θ and ϕ from uniform distributions. Each element of θ is sampled from a uniform distribution from 0 to a fixed fraction of the size of the image in that dimension. Each element of ϕ is sampled from a uniform distribution of 1 to $2k + 1$ for some k . We have not been able to tune parameters much, and currently the parameters are set at $\frac{1}{6}$ of the image size for θ , and $k = 5$ for ϕ .

4 RESULTS

Our current results are limited by the fact that the code takes a long time to run on a laptop. We are currently trying to get the code running on machines with more computing power, but there are some compiling issues on Linux that have not yet been resolved. However, running the data on very small random forests (1 tree with a depth of 2) with very sparse sampling (one voxel for every 7x7x3 box), we get that the predicted bounding box of the test patient is within 10% of the actual bounding box.

[I didn't think it was worth writing up my very preliminary results for this draft paper if I would have to change it for the final anyways.]

5 BIBLIOGRAPHY

[1] D. Pace, A. Dalca, T. Geva, A. Powell, M. Moghari, P. Golland, Interactive Whole-Heart Segmentation in Congenital Heart Disease

[2] A. Criminisi, J. Shotton (eds.), Decision Forests for Computer Vision and Medical Image Analysis, Advances in Computer Vision and Pattern Recognition